

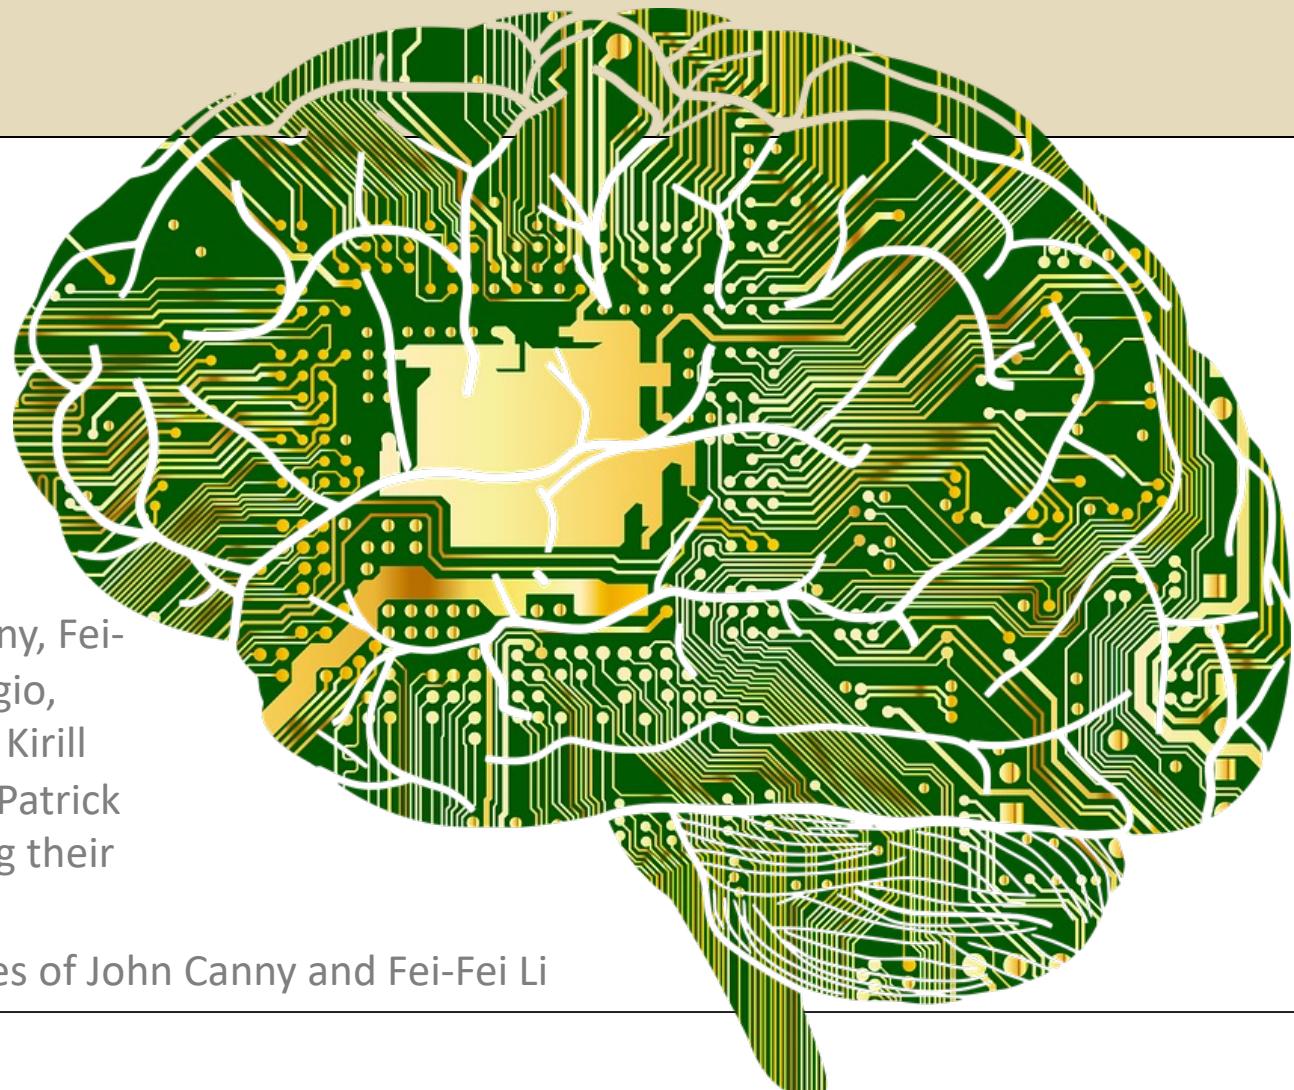
Deep Learning

Architectures and Methods

Introduction



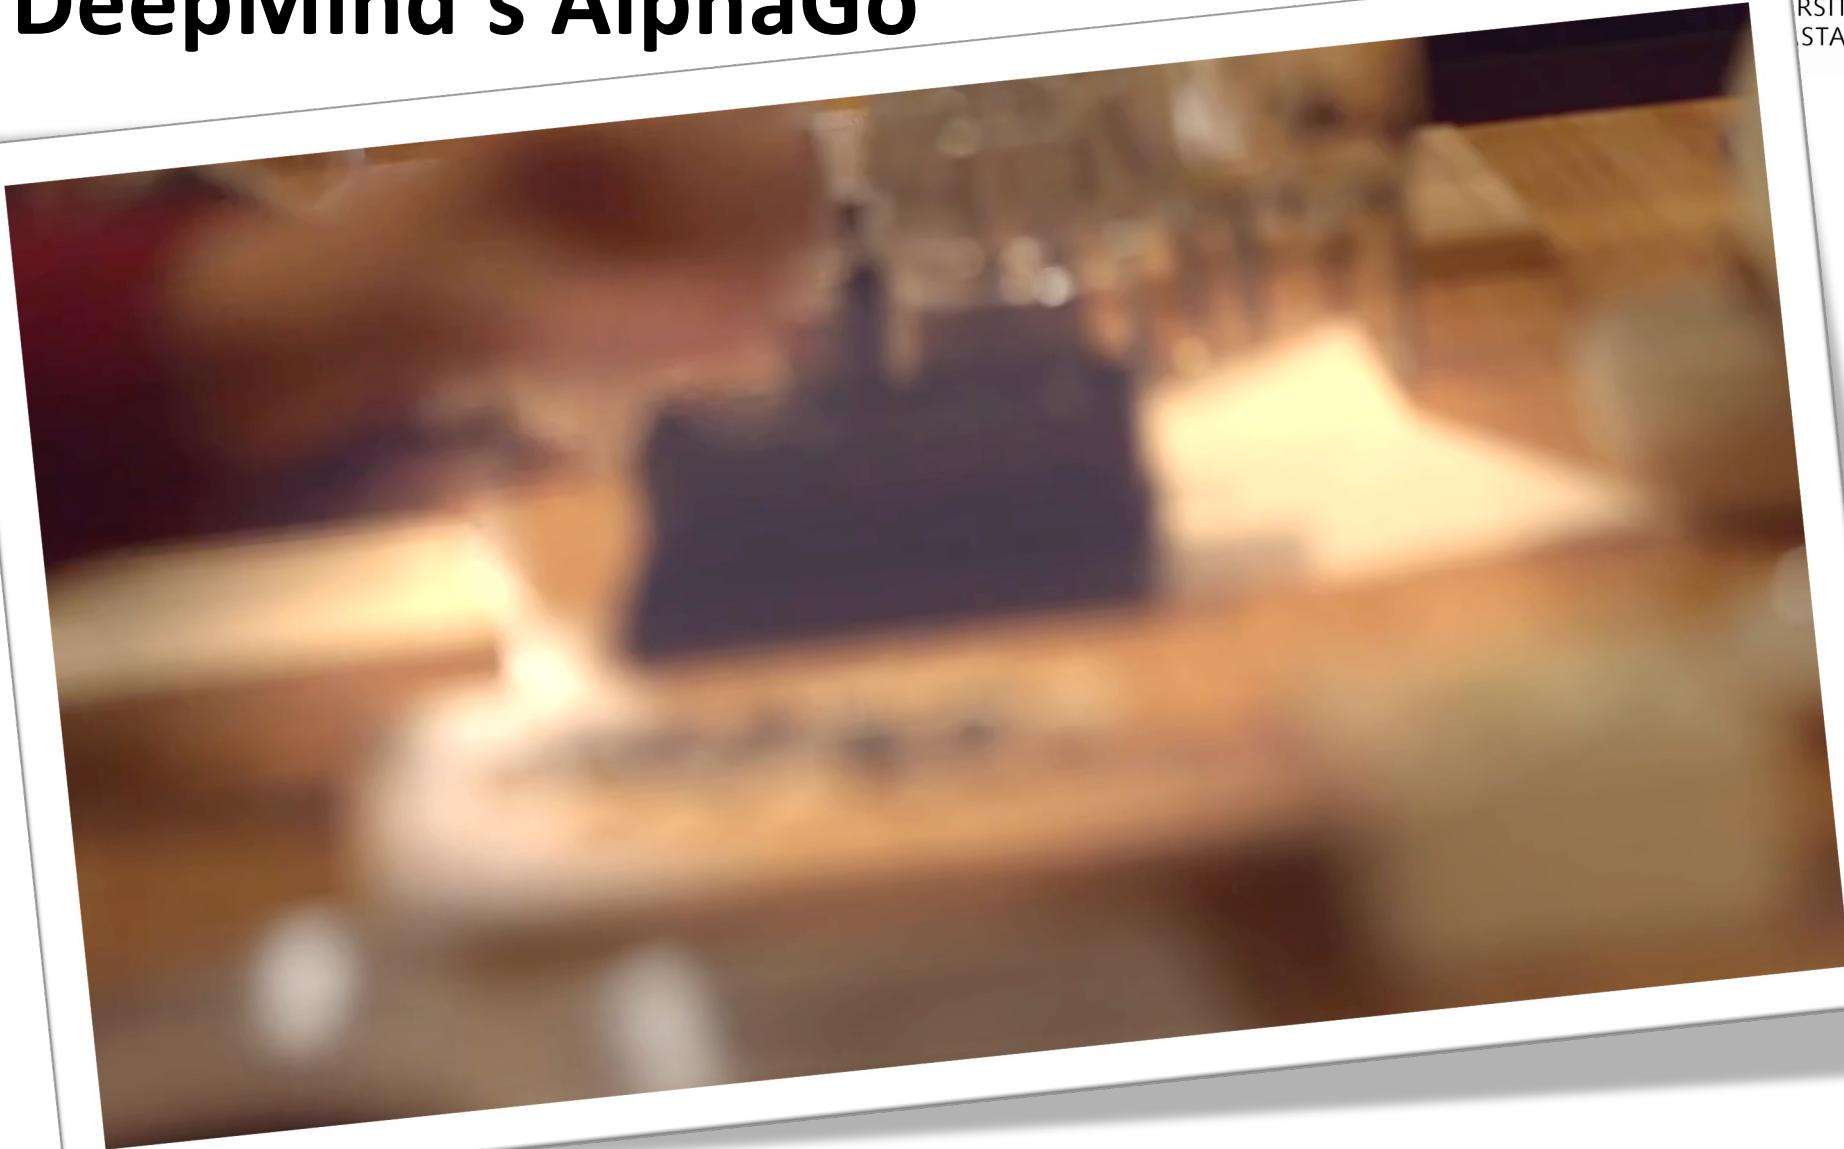
TECHNISCHE
UNIVERSITÄT
DARMSTADT



Thanks to Sergey Levine, John Canny, Fei-Fei Li, Ian Goodfellow, Yoshua Bengio, Aaron Courville, Efstratios Gavves, Kirill Gavrilyuk, Berkay Kicanaoglu, and Patrick Putzky and many others for making their materials publically available.

The slides are mainly based on ones of John Canny and Fei-Fei Li

DeepMind's AlphaGo



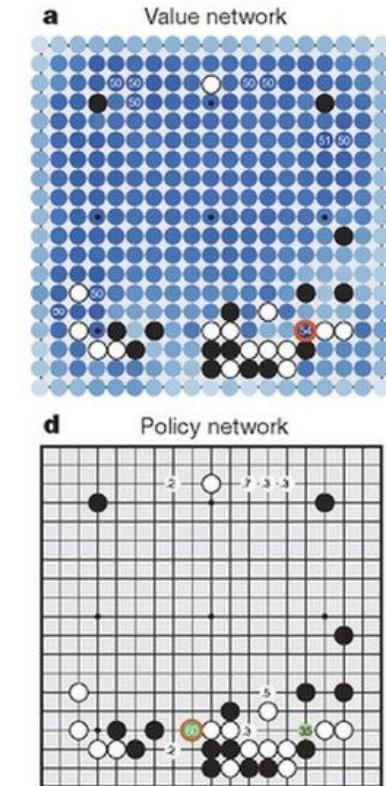
Watch NATURE video at <https://www.youtube.com/watch?v=g-dKXOlsf98>



DeepMind's AlphaGo



TECHNISCHE
UNIVERSITÄT
DARMSTADT



Deep policy network is trained to produce probability map of promising moves. The deep value network is used to prune the search tree (monte-carlo tree search); so there is a lot of classical AI machinery around the deep part.



And yes, the machine may also learn to play other games



Goal of Deep Architectures

To this aim most approaches use (stacked) neural networks

High-level semenatical representations

Edges, local shapes, object parts

Low level representation

Deep learning methods aim at

- **learning feature hierarchies**
- where features from higher levels of the hierarchy are formed by lower level features.

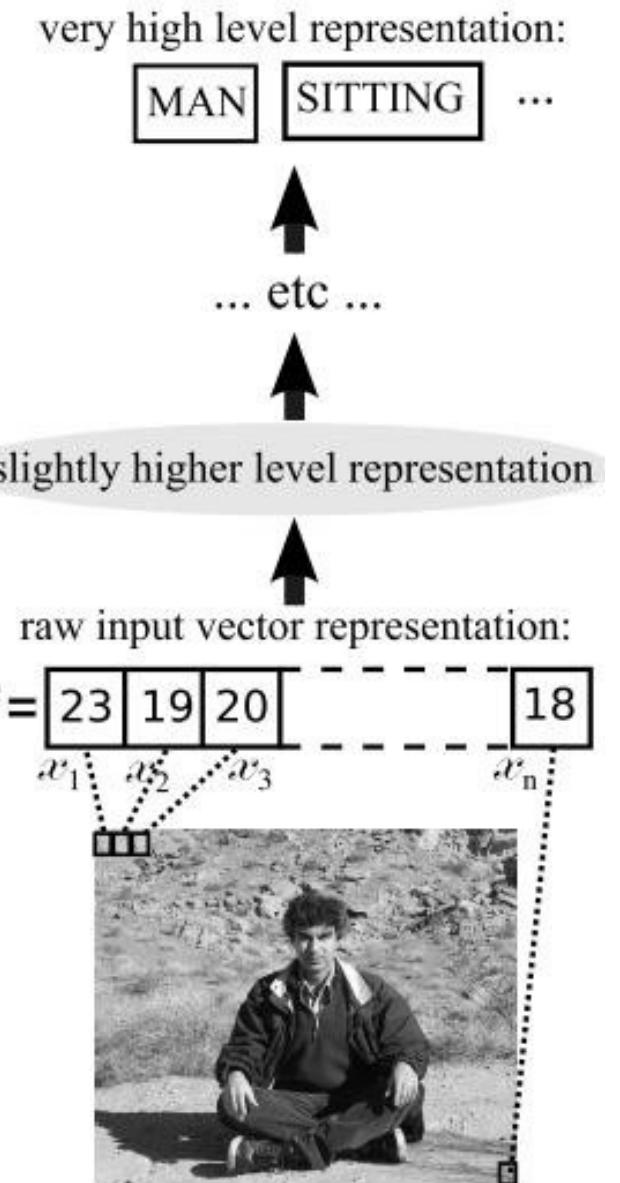


Figure is from Yoshua Bengio

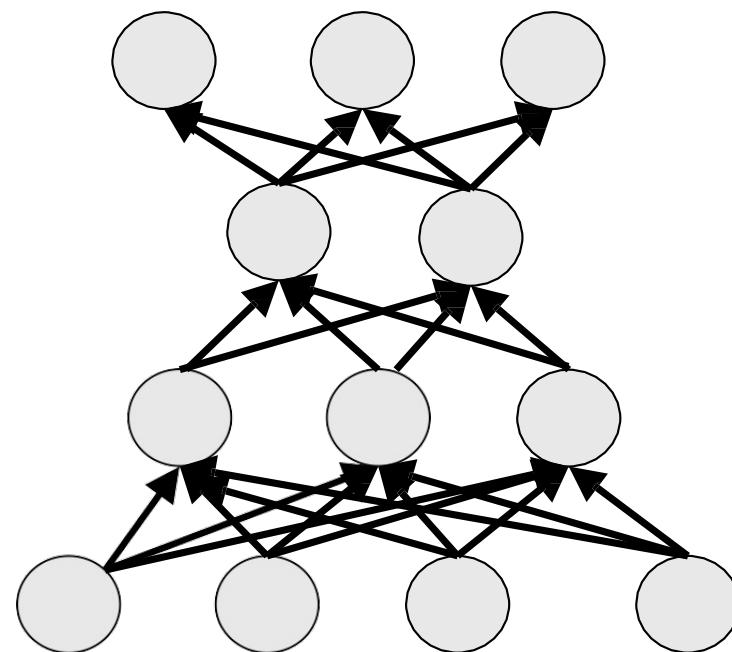
Deep Architectures

Deep architectures are composed of multiple levels of non-linear operations, such as neural nets with many hidden layers.

Output layer

Hidden layers

Input layer



Examples of non-linear activations:

$$\tanh(x)$$

$$\sigma(x) = (1 + e^{-x})^{-1}$$

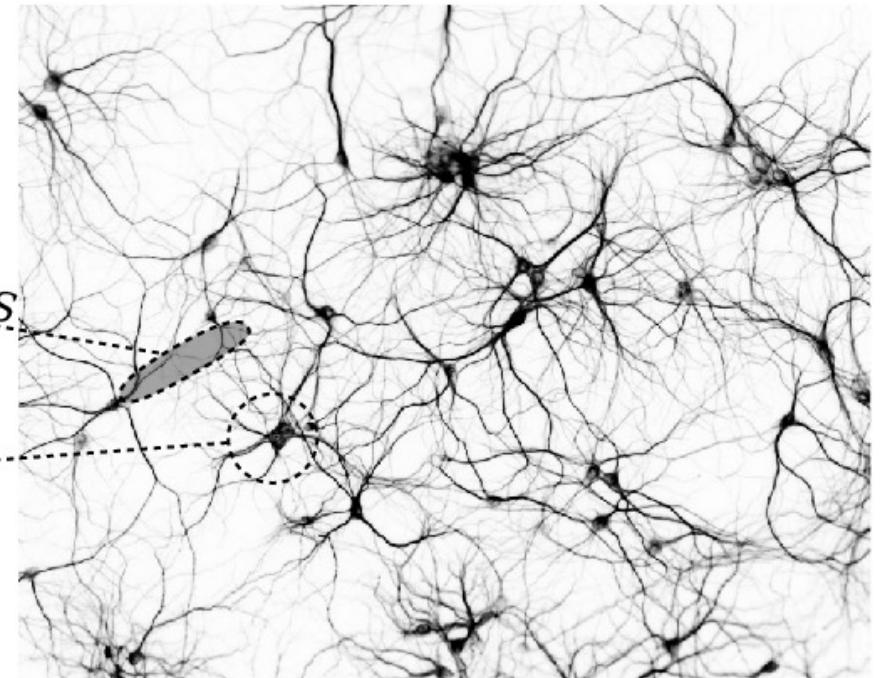
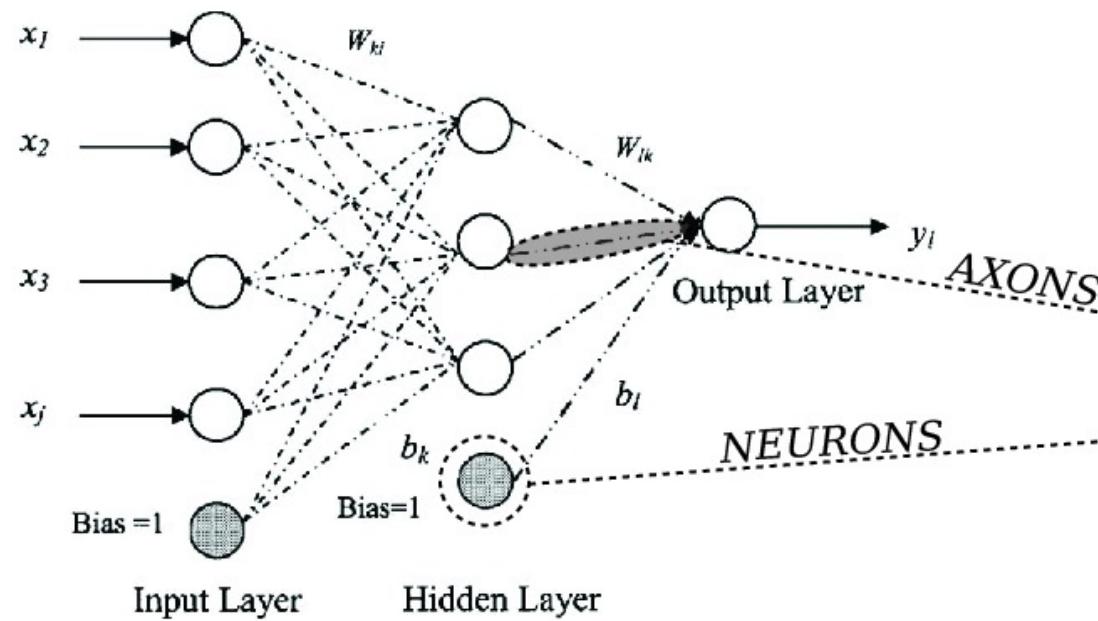
$$\max(0, x)$$

In practice, NN with multiple hidden layers work better than with a single hidden layer.



Artificial Neural Networks are inspired by neural networks

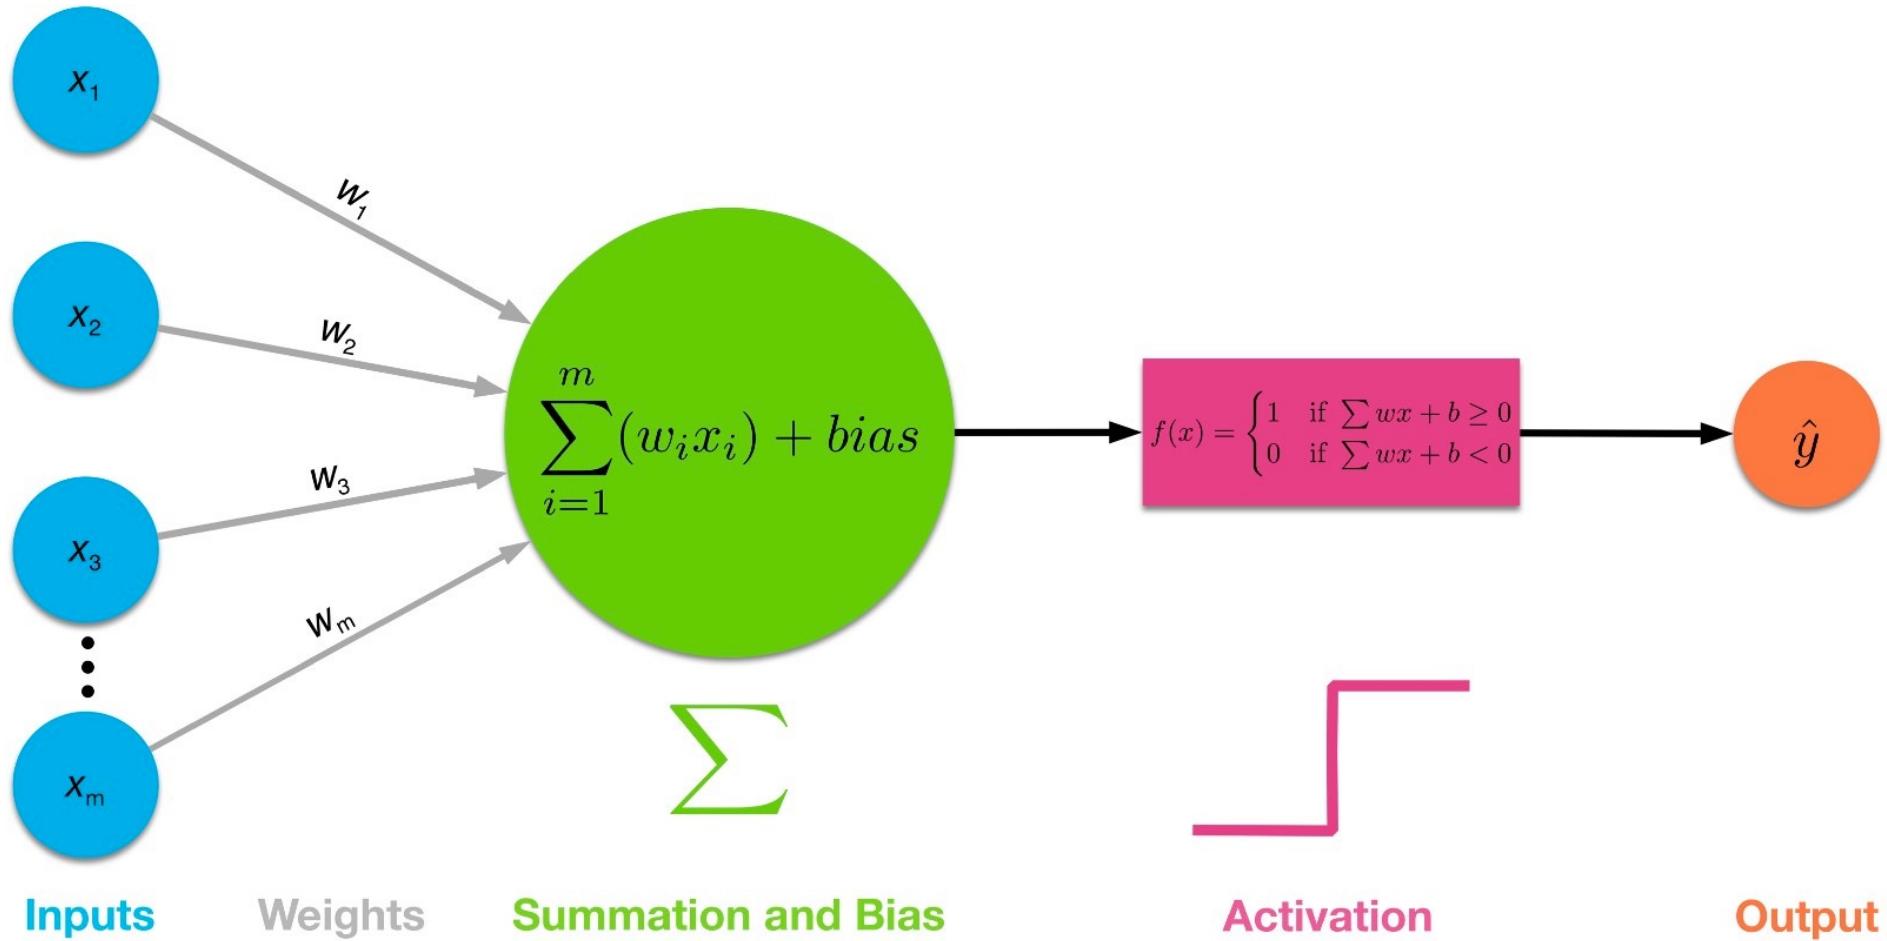
NEURAL NETWORK MAPPING



Abstract Neural Unit



TECHNISCHE
UNIVERSITÄT
DARMSTADT





Godzillium vs. Trumpium:
Some Suggestions to Add
to the Periodic Table



To Protect Against Zika
Virus, Pregnant Women
Are Warned About Latin
American Trips



THE NEW OLD
F.T.C.'s Lum
Doesn't End
Training Del

nature

International weekly journal of science

SCIENCE

And this has produced a lot of media echo

Scientists See Promise in Deep-Learning Program

By JOHN MARKOFF NOV. 23, 2012



NEWS

Sign in

News

Sport

Weather

Show

Home

Video

World

UK

Business

Tech

Science

Magazi

NATURE | NEWS

عربى

Game-playing software holds lessons for neuroscience

DeepMind computer provides new way to investigate how the brain

Forbes / Tech

DEC 29, 2014 @ 11:37 AM 89,471 VIEWS

Tech 2015: Deep Learning And Machine Intelligence Will Eat The World



'Deep learning' technology inspired by human brain

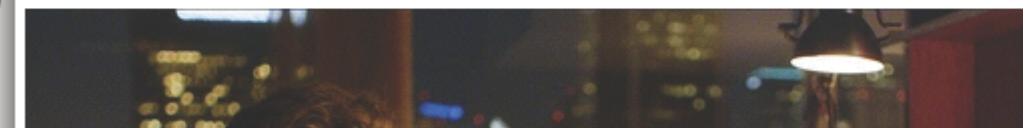
culture business lifestyle fashion environment tech travel

Droids do dream of electric sheep

in feedback loop in its image recognition neural network - which

Google a step closer to developing machines with human-like intell

Algorithms developed by Google designed to encode thoughts, could computers with 'common sense' within a decade, says leading AI



RCH

h



to

SCIENCE



Scientis

By JOHN MARKOFF



Forbes

DEC 29, 2014 @ 11:37 AM

Tech 20
Eat The



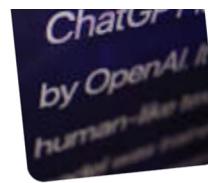
ndroids do

in feedback loop in its

E EL PAÍS in English

ChatGPT: How to use it and what can it really do

Anyone can use this surprising and impressive artificial intelligence tool, created by OpenAI.
vor 3 Wochen



ekly journal of science

echo



holds lessons

investigate how the bra



stocks for 2016



velopin
ke intell

e thoughts, co
ys leading AI



The Guardian

ChatGPT allowed in International Baccalaureate essays

Content created by chatbot must be treated like any other source and attributed when used, says IB.

vor 1 Monat

Kyodo News

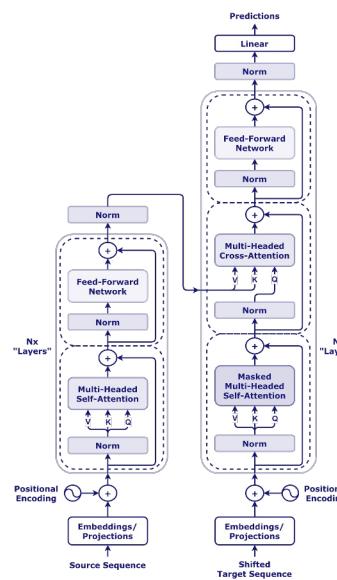
ChatGPT developer eager to open office in Japan amid



It is all about attention and scale

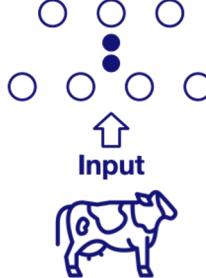


TECHNISCHE
UNIVERSITÄT
DARMSTADT

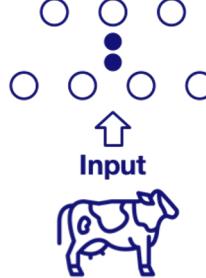


Supervised
implausible labels

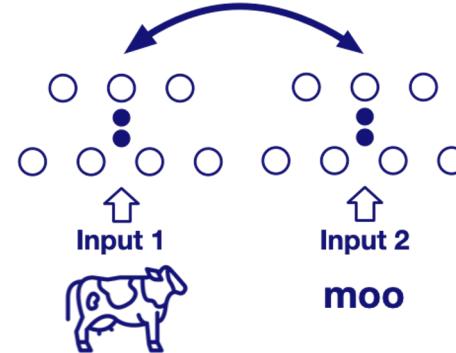
"COW"
Target



Unsupervised
limited power



Self-supervised
derives label from a
co-occurring input to
related information



Transformer

Self-Supervised Learning

Scale

AI Research Director at Deepmind says all we need now is scaling



Nando de Freitas @Nando... · 4 t.
Someone's opinion article. My opinion:
It's all about scale now! The Game is
Over! It's about making these models
bigger, safer, compute efficient, faster at
sampling, smarter memory, more
modalities, INNOVATIVE DATA, on/
offline, ... 1/N



thenextweb.com
DeepMind's new Gato AI makes me
fear humans will never achieve AGI

10 22 78 1



Google Engineer Says Chat Bot Is Sentient

Bloom

June 24, 2022

<https://www.youtube.com/watch?v=kgCUn4fQTsc>



But this might not be the human way



TECHNISCHE
UNIVERSITÄT
DARMSTADT

Current Biology

Explore Online Now Current Issue Archive Journal Information For Authors

< Previous Article Volume 27, Issue 18, p2827–2832.e3, 25 September 2017 Next Article >

REPORT

Humans, but Not Deep Neural Networks, Often Miss Giant Targets in Scenes

Miguel P. Eckstein⁴, Kathryn Koehler, Lauren E. Welbourne, Emre Akbas

⁴ Lead Contact

Summary

Even with great advances in machine vision, animals are still unmatched in their ability to visually search complex scenes. Animals from bees [1, 2] to birds [3] to humans [4, 5, 6, 7, 8, 9, 10, 11, 12] learn about the statistical relations in visual environments to guide and aid their search for targets. Here, we investigate a novel manner in which humans utilize rapidly acquired information about scenes by guiding search toward likely target sizes. We show that humans often miss targets when their size is inconsistent with the rest of the scene, even when the targets were made larger and more salient and observers fixated the target. In contrast, we show that state-of-the-art deep neural networks do not exhibit such deficits in finding mis-scaled targets but, unlike humans, can be fooled by target-shaped distractors that are inconsistent with the expected target's size within the scene. Thus, it is not a human deficiency to miss targets when they are inconsistent in size with the scene; instead, it is a byproduct of a useful strategy that the brain has implemented to rapidly discount potential distractors.



Anyhow, deep neural learning fueled the discussion on AI



Stephen Hawking

"Success in creating AI would be the biggest event in human history...."

"Unfortunately, it might also be the last, unless we learn how to avoid the risks. In the near term, world militaries are considering autonomous-weapon systems that can choose and eliminate targets."

"...humans, limited by slow biological evolution, couldn't compete and would be superseded by A.I."

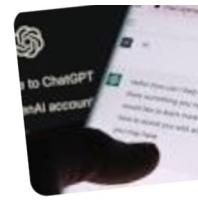




An Open Letter Asks AI Researchers To Reconsider Responsibilities

In recent months, it's been hard to escape hearing about artificial intelligence platforms such as ChatGPT, the AI-enabled version of Bing,...

vor 5 Tagen



d the on AI

'AI Pause' Open Letter Stokes Fear and Controversy

The recent call for a six-month "AI pause"—in the form of an online letter demanding a temporary artificial intelligence moratorium—has...

vor 4 Tagen



Open Letter Urges Pause on AI Research

The Future of Life Institute has issued an open letter calling for a six-month pause on some forms of AI research. Citing "profound risks to..."

vor 2 Wochen



Tech 24 - AI open letter authors: 'We want people to share our level of uncertainty'

The authors of an open letter advocating a "pause" in the development of advanced artificial intelligence have urged researchers to back...

vor 2 Wochen



1,100+ notable signatories just signed an open letter asking all AI labs to immediately pause for at least 6 months'

Elon Musk and Stephen Hawking are among the signatories.



Artificial Intelligence and Humans

"How do we humans get so much from so little?" and by that I mean how do we acquire our understanding of the world given what is clearly by today's engineering standards so little data, so little time, and so little energy.



Josh Tenenbaum
“Bayesian Program Learning”



Lake, Salakhutdinov, Tenenbaum,
Science 350 (6266), 1332-1338, 2015
Tenenbaum, Kemp, Griffiths, Goodman,
Science 331 (6022), 1279-1285, 2011

Systems AI: the computational and mathematical modeling of complex AI systems.



OpenAI



Deep Learning: Hype or Hope?



TECHNISCHE
UNIVERSITÄT
DARMSTADT

Hype: “extravagant or intensive publicity or promotion”

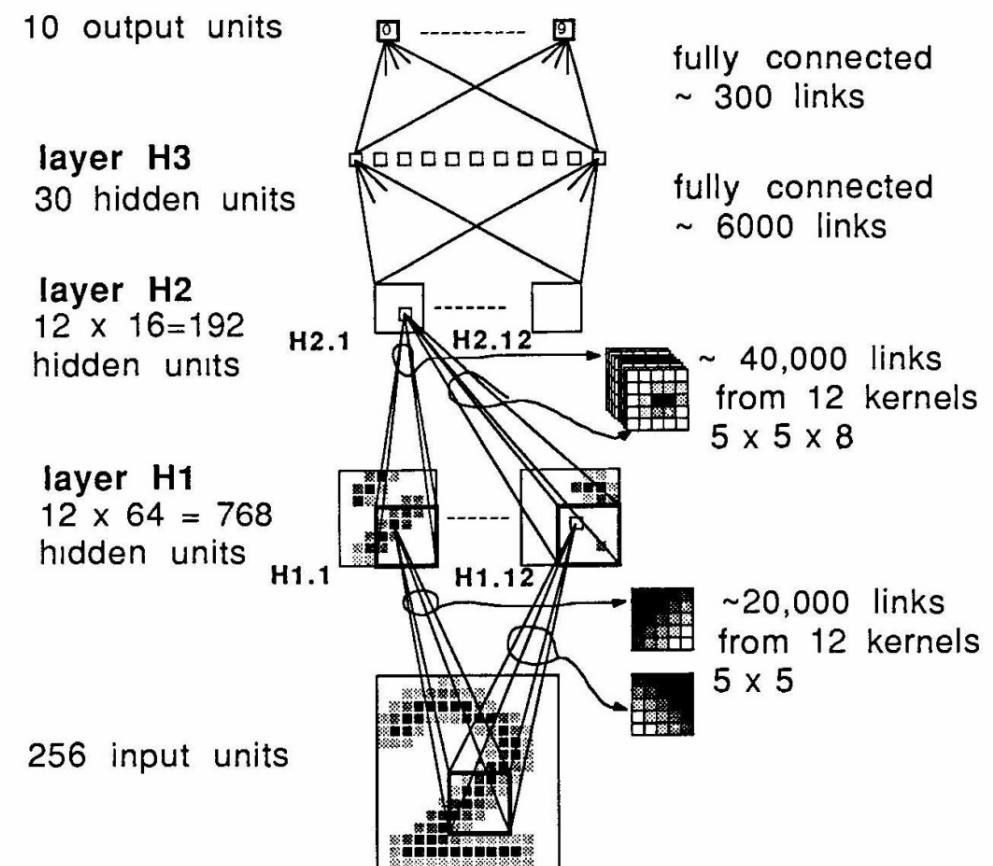
Hope: “expectation of fulfillment or success”



Milestones: Digit Recognition

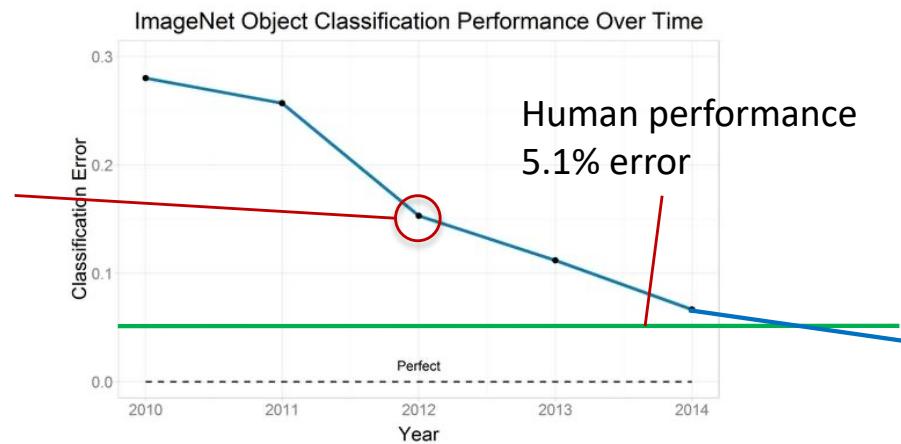
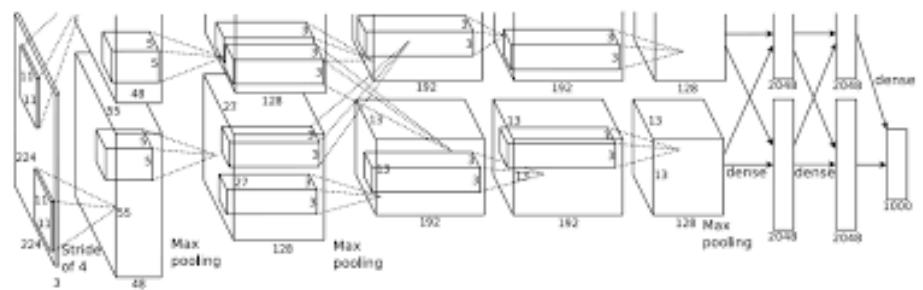
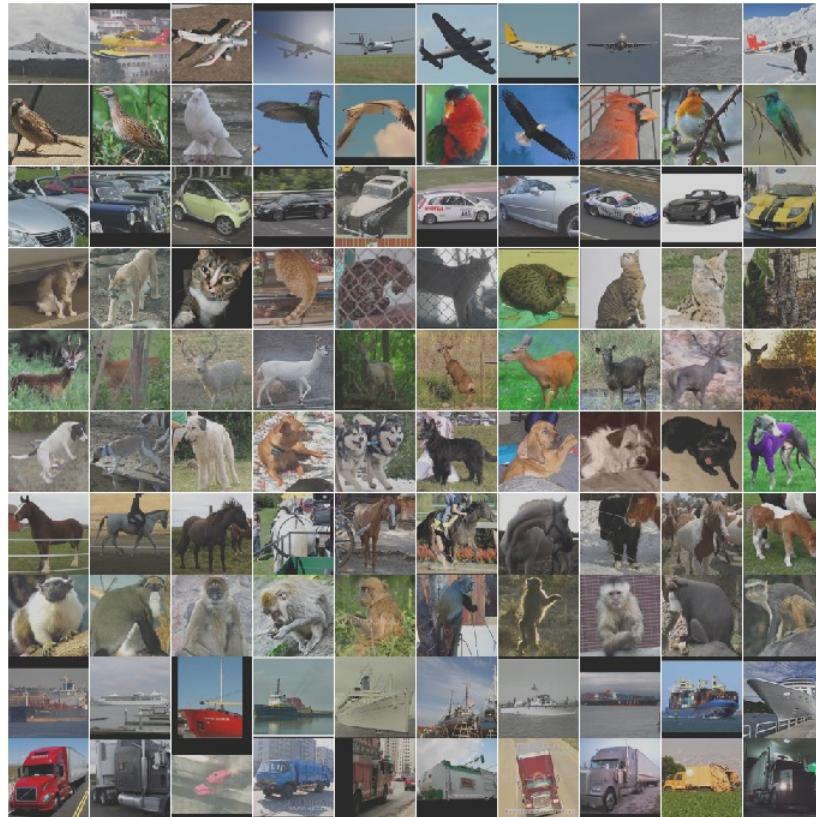
LeNet 1989: recognize zip codes, Yann Lecun, Bernhard Boser and others, ran live in US postal service

80322 - 4129 80206
40004 14310
37872 05153
35502 75216
35460 44209



Milestones: Image Classification

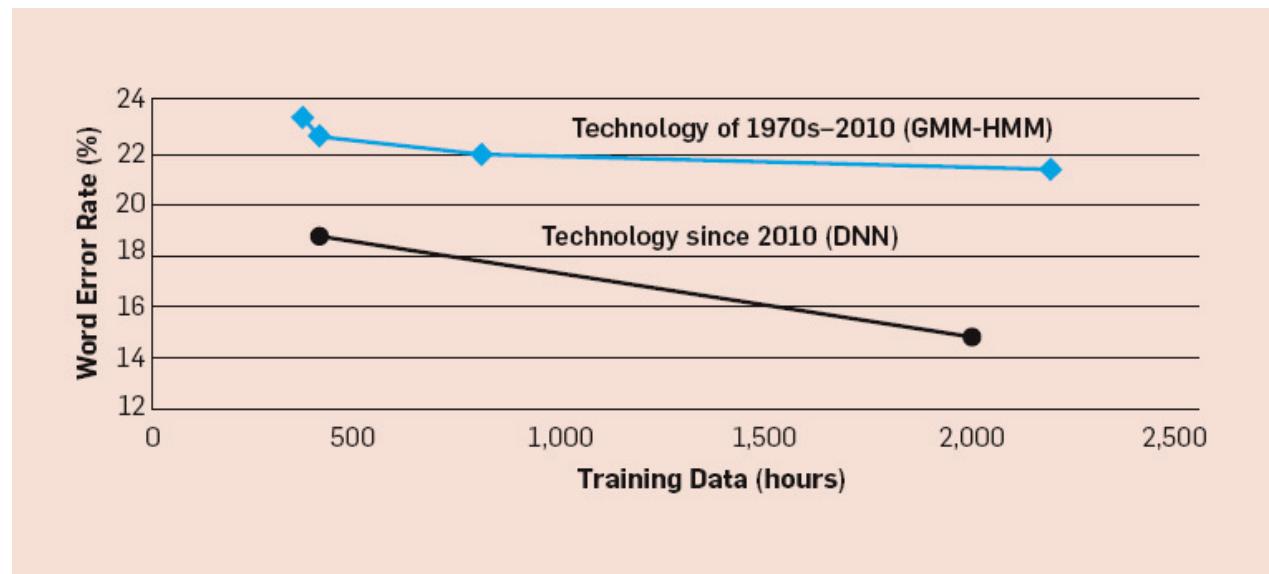
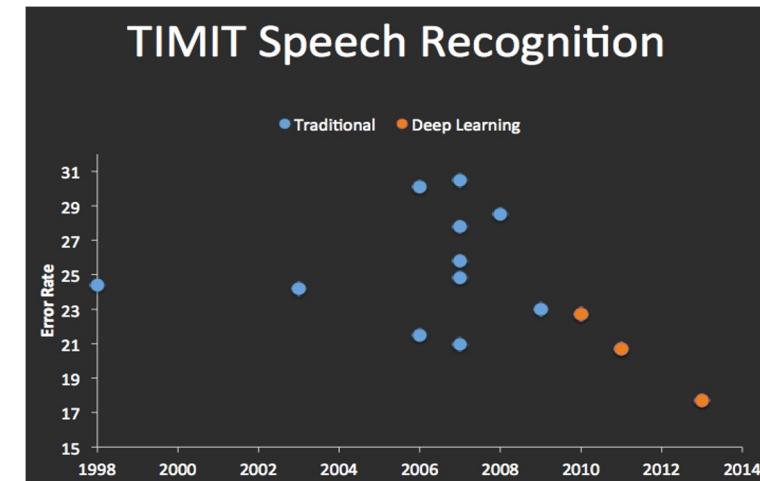
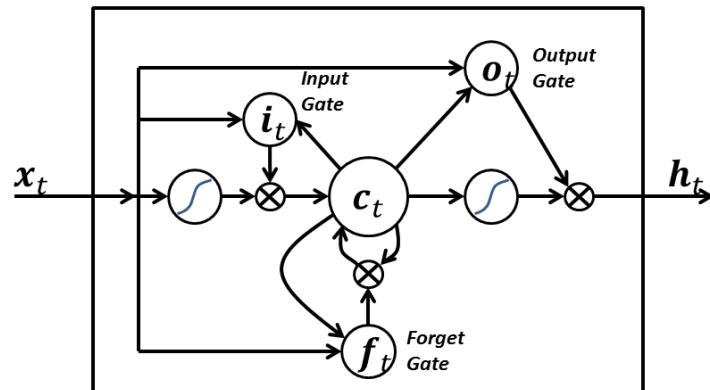
Convolutional NNs: AlexNet (2012): trained on 200 GB of ImageNet Data



Milestones: Speech Recognition



Recurrent Nets: LSTMs (1997):

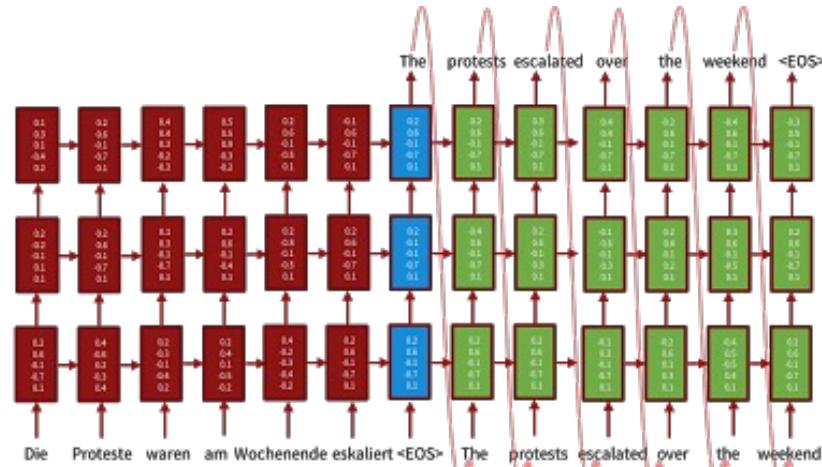


Milestones: Language Translation



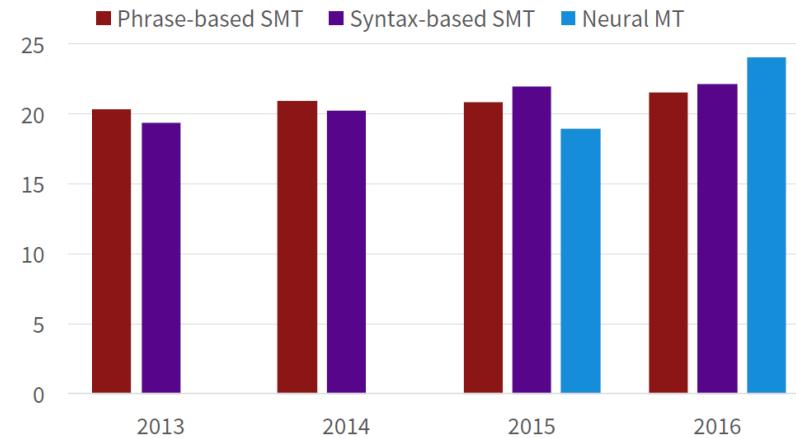
TECHNISCHE
UNIVERSITÄT
DARMSTADT

Sequence-to-sequence models with LSTMs and attention:



Progress in Machine Translation

[Edinburgh En-De WMT newstest2013 Cased BLEU; NMT 2015 from U. Montréal]



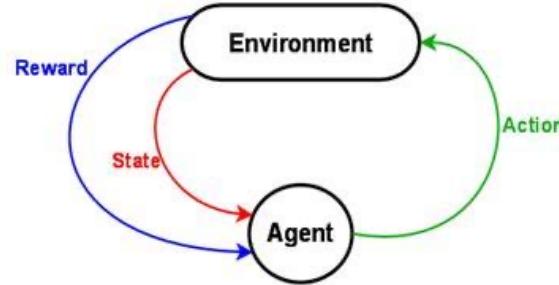
From [Sennrich 2016, http://www.meta-net.eu/events/meta-forum-2016/slides/09_sennrich.pdf]

Source Luong, Cho, Manning ACL Tutorial 2016.

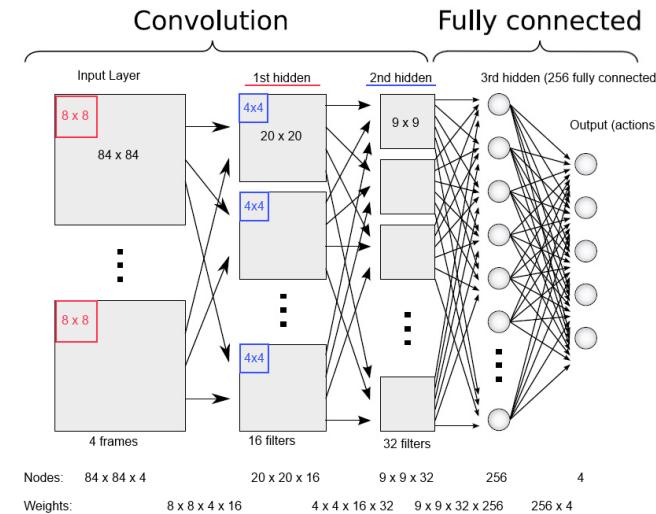


Milestones: Deep Reinforcement Learning

In 2013, Deep Mind's arcade player bests human expert on six Atari Games. Acquired by Google in 2014.



In 2016, Deep Mind's alphaGo defeats former world champion Lee Sedol

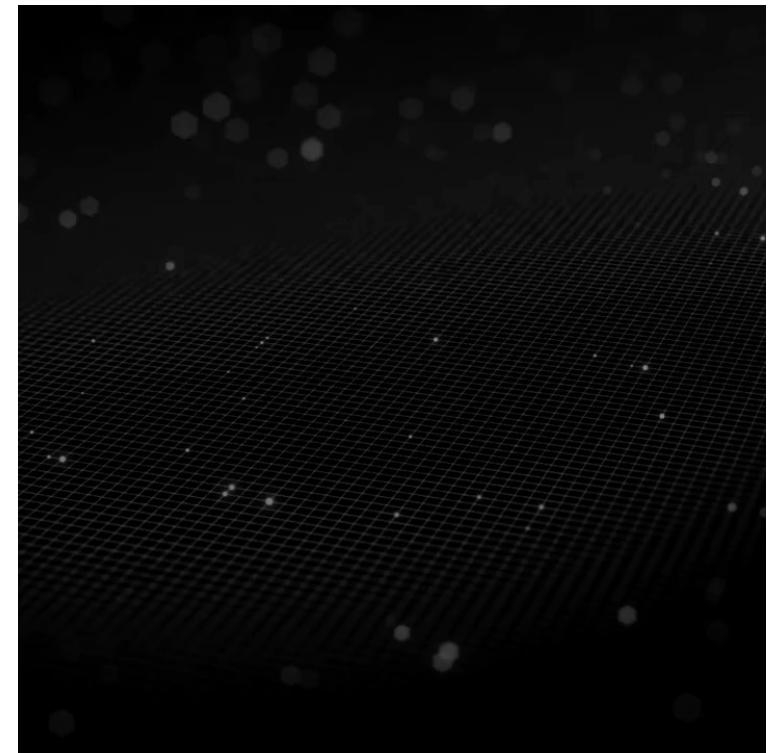


Milestones: ChatGPT, GPT-4, ... here Luminos



Aleph Alpha

Disclaimer: I am an investor of Aleph Alpha



<https://www.youtube.com/channel/UCx6PVRx2wRVbaYHQC7-axtg/videos>



Deep Learning: Is it Hype or Hope?



TECHNISCHE
UNIVERSITÄT
DARMSTADT



Deep Learning: Is it Hype or Hope?



TECHNISCHE
UNIVERSITÄT
DARMSTADT

Yes!

But ...



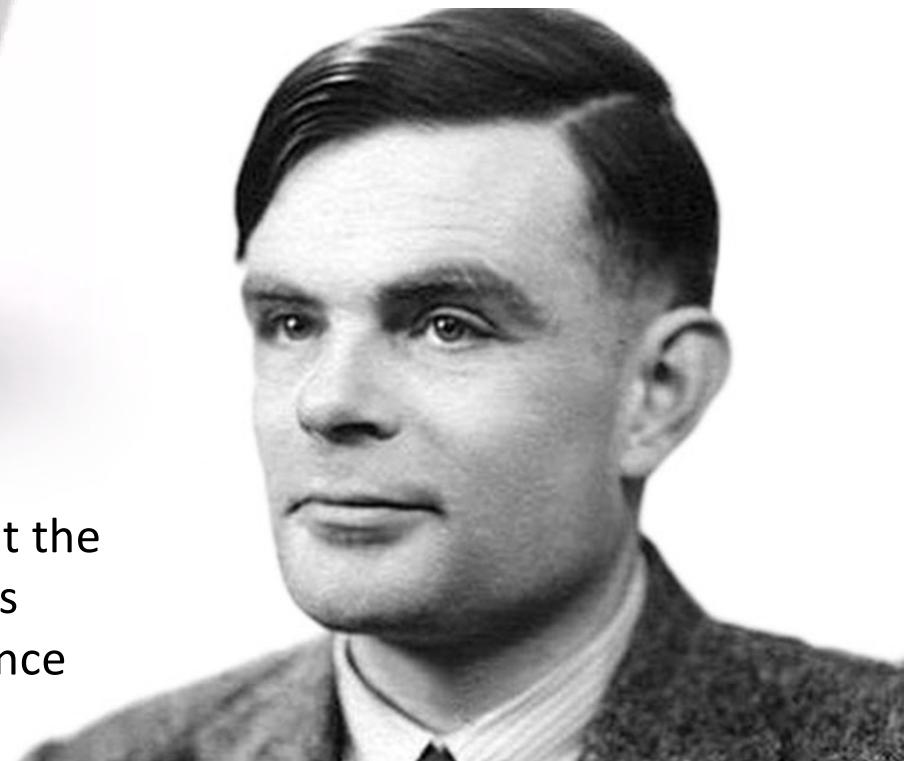
Deep Learning honored with Turing Award 2019



Yann LeCun, Geoffrey Hinton, and Yoshua Bengio



Turing Award = Nobel Prize for Computing



Named after Alan Turing, a British mathematician at the University of Manchester. Turing is often credited as being the key founder of theoretical computer science and AI.



Learning about Deep Neural Networks

Yann Lecun quote: DNNs require: “an interplay between intuitive insights, theoretical modeling, practical implementations, empirical studies, and scientific analyses”



I.e. there isn't a framework or core set of principles to explain everything (c.f. graphical models for machine learning) yet. **We try to cover the ground in Lecun's quote.**



This Course

Goals:

- Introduce deep learning to a broad audience.
- Review principles and techniques for understanding deep networks.
- Develop skill at designing networks for applications.

There are several other DL courses at TU Darmstadt that often focus on specific domains. Check them out!



This Course: a partly flipped classroom

We will mainly follow Fei-Fei Li's (Stanford) course **CS231n: Convolutional NNs for Visual Recognition**

<http://cs231n.stanford.edu/>

and Sergey Levine's (UC Berkeley) course **CS294-129 Designing, Visualizing & Understanding Deep NNs**

<https://bcourses.berkeley.edu/courses/1453965/pages/cs294-129-designing-visualizing-and-understanding-deep-neural-networks>

and own ones from previous years!

This will be done as a partly flipped classroom experiment: **videos are used to 'deliver content' outside of the classroom. Online lecture to recap important parts. Then questions from you**



This Course: a partly flipped classroom

You will watch the next video lecture prior to the lecture, **prepare a list of 5 questions as main homework for that lecture**, and then we will try to **answer your question in class as a group**.

- We may also present our own stuff without videos. We may even ask you to read papers
- Next to the “5-questions” homework, there are also textual exercises and you can propose and run a deep learning project.



This Course



We may go for the bonus option. If so, roughly:

- Class Participation: 30%
- Questions: 30%
Let's see how
we will do this.
Not decided yet
- Final Project (in groups): 40%

Main Audience: CS students

Final exam will be a written exam



Logistics



TECHNISCHE
UNIVERSITÄT
DARMSTADT

- Instructor: Kersting
- Time: Monday 11:30-13:10, S306/051

- Exercise: Wednesday 11:40-13:20, S202/C205 Bosch
- Tutors: Hikaru Shindo, Jonas Seng



Course Project



TECHNISCHE
UNIVERSITÄT
DARMSTADT

- Decision and more info later
- 2-pages expose to get a deal on a project
- Encourage “open-source” projects that can be archived somewhere.
- Talk also to other groups at TUDa



Tentative Outline



1. Intro, Computer Vision History, Classification, KNN
2. Linear Classification, Feature Selection,
Optimization, Stochastic Gradient,
Backpropagation
3. Training DNNs: Activation functions, initialization,
gradient flow, batch normalization, parameter
updates, ensembles, dropout
4. Convolutional Neural Networks
5. Recurrent Networks, LSTMs



Tentative Outline



TECHNISCHE
UNIVERSITÄT
DARMSTADT

6. Deep NLP
7. Variational Autoencoder
8. Deep Probabilistic Models
9. Interpreting DNNs
10. DNNs for Games
11. Deep Reinforcement Learning
12. Final Project Presentations
13. Final Project Presentations

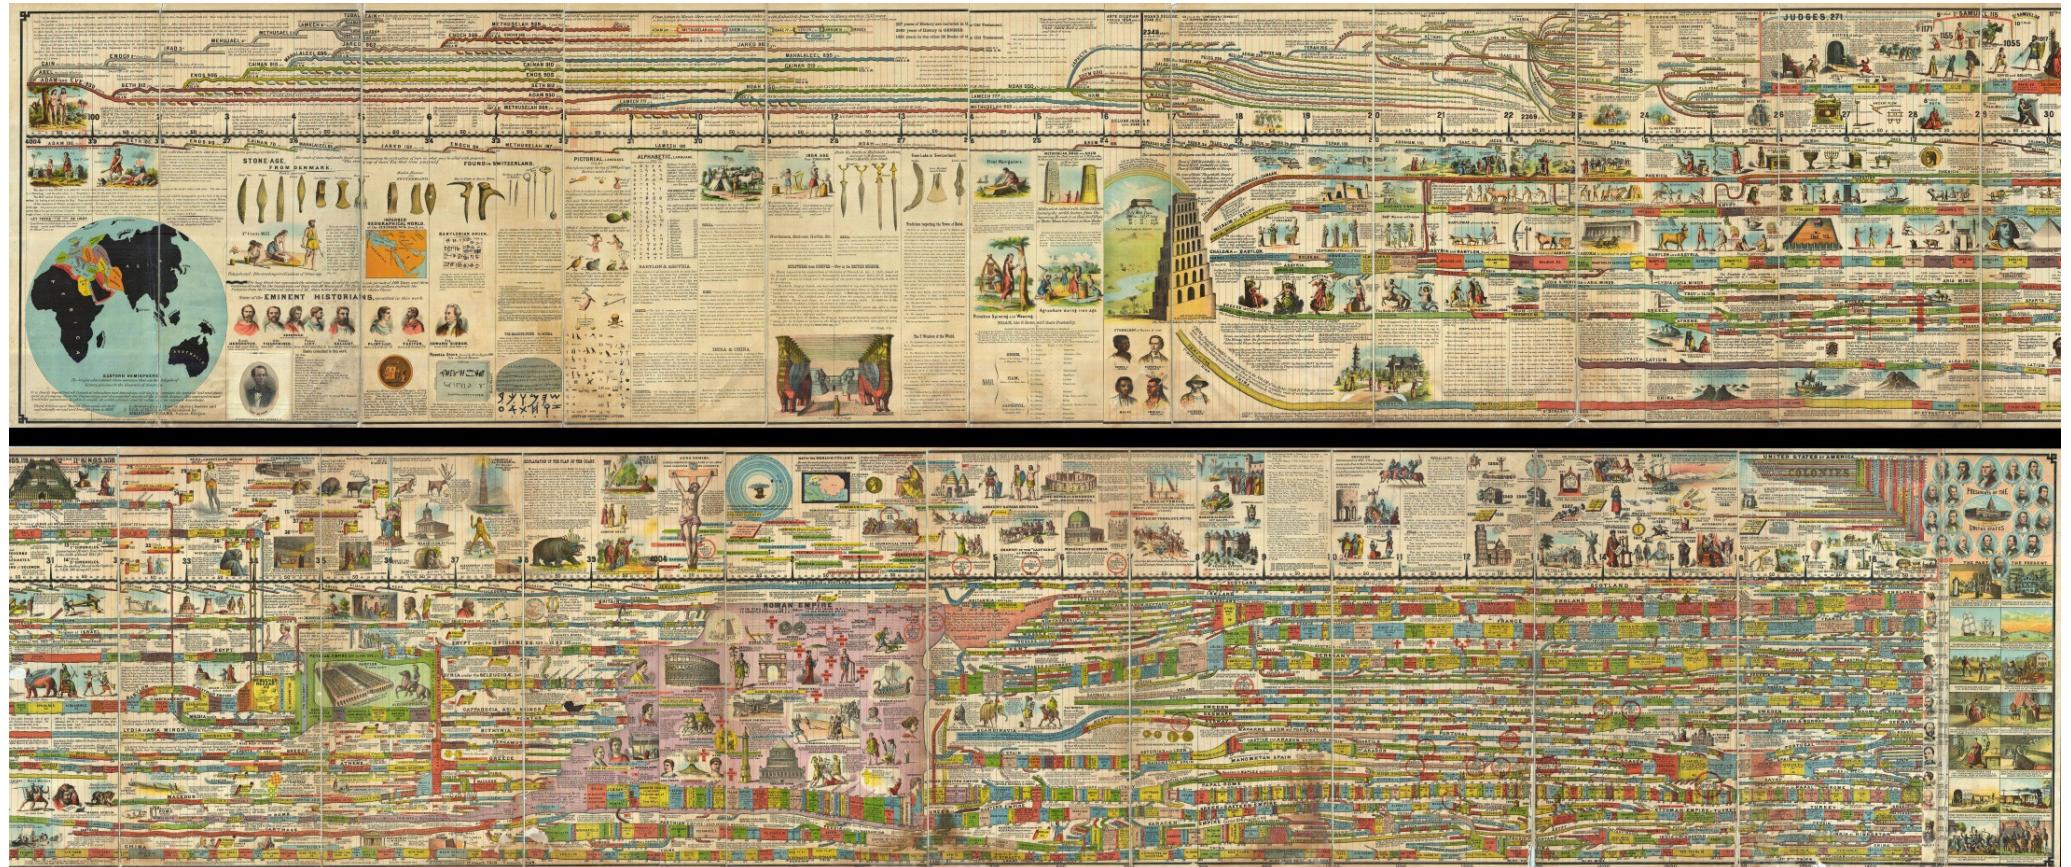


Some History



TECHNISCHE
UNIVERSITÄT
DARMSTADT

- Reading: the Deep Learning Book, Introduction



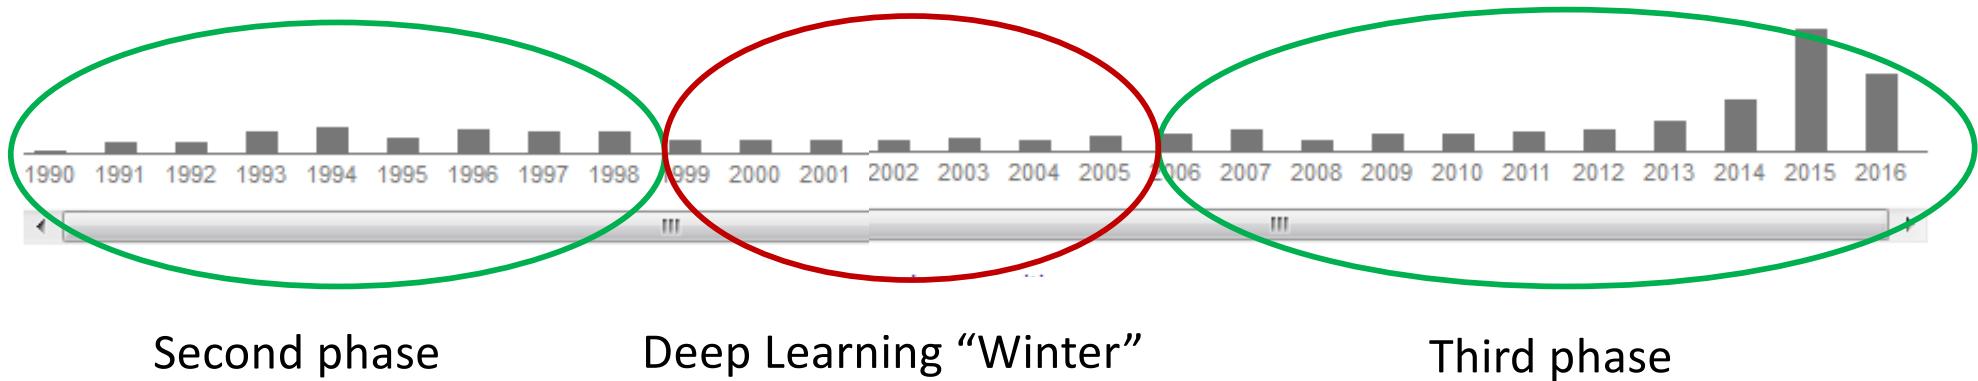
Phases of Neural Network Research

- 1940s-1960s: Cybernetics: Brain like electronic systems, morphed into modern control theory and signal processing.
- 1960s-1980s: Digital computers, automata theory, computational complexity theory: simple shallow circuits are very limited...
- 1980s-1990s: Connectionism: complex, non-linear networks, back-propagation.
- 1990s-2010s: Computational learning theory, graphical models: Learning is computationally hard, simple shallow circuits are very limited...
- 2006→: Deep learning: End-to-end training, large datasets, explosion in applications.



Citations of the “LeNet” paper

- Recall the LeNet was a modern visual classification network that recognized digits for zip codes. Its citations look like this:

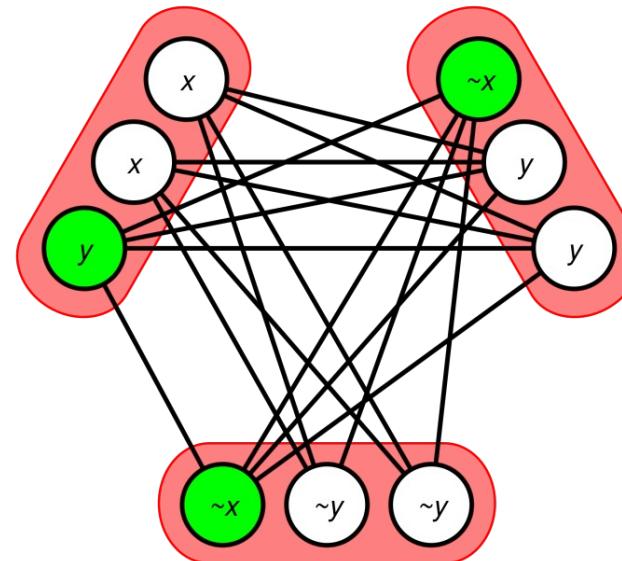


- The 2000s were a golden age for machine learning, and marked the ascent of graphical models. But not so for neural networks.



Why the success of DNNs is surprising

- From both complexity and learning theory perspectives, simple networks are very limited.
 - Can't compute parity with a small network.
 - NP-Hard to learn “simple” functions like 3SAT formulae, and i.e. training a DNN is NP-hard.



Why the success of DNNs is surprising

- The most successful DNN training algorithm is a version of gradient descent which will only find local optima. In other words, it's a greedy algorithm. Backprop:

$$\text{loss} = f(g(h(y)))$$

$$d \text{ loss}/dy = f'(g) \times g'(h) \times h'(y)$$

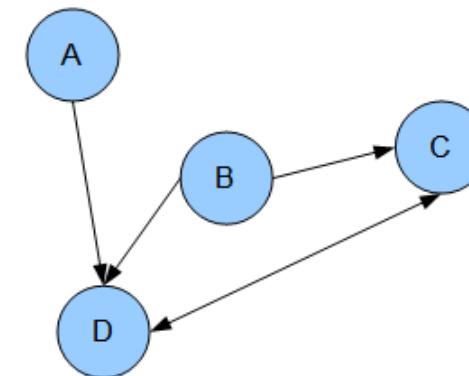
- Greedy algorithms are even more limited in what they can represent and how well they learn.
- If a problem has a greedy solution, its regarded as an “easy” problem.



Why the success of DNNs is surprising

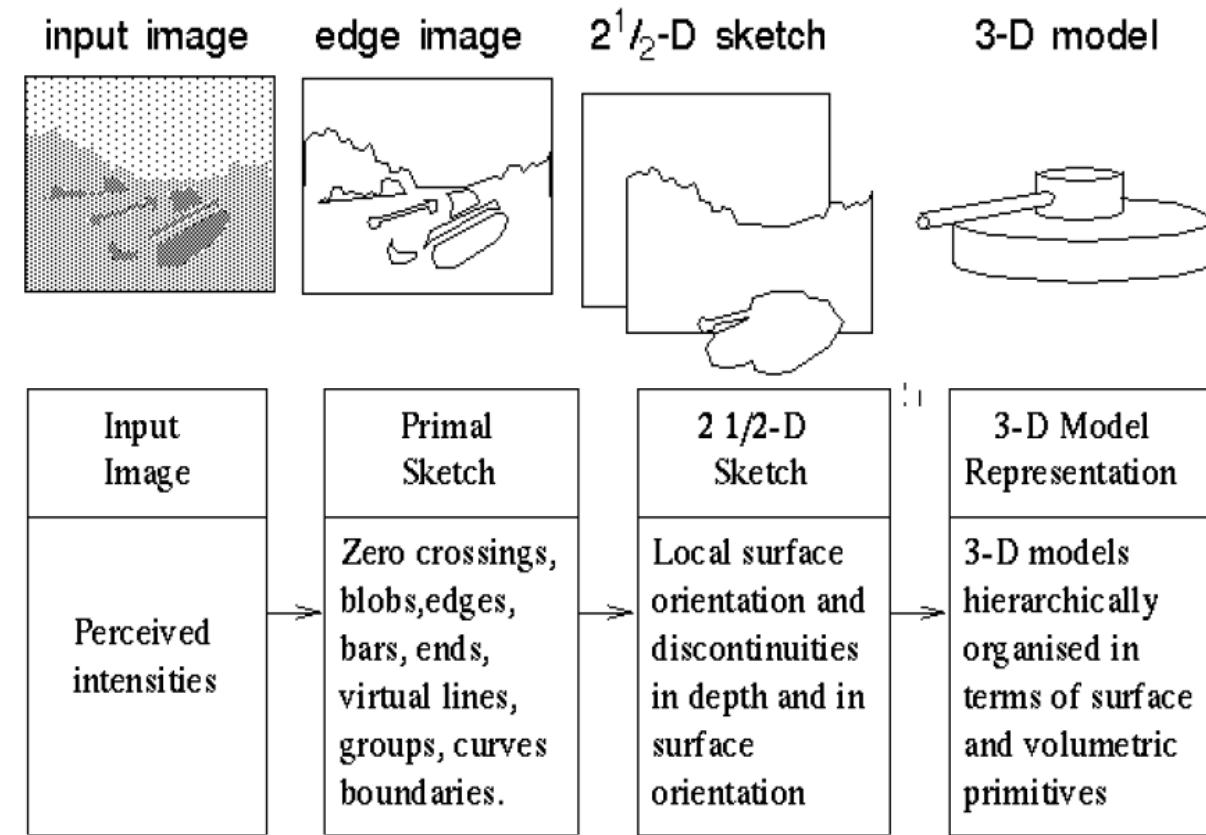
- In graphical models, values in a network represent random variables, and have a clear meaning. The network structure encodes dependency information, i.e. you can represent rich models.
- In a DNN, node activations encode nothing in particular, and the network structure only encodes (trivially) how they derive from each other.

**But more on the connection
between DNNs and PGMs
later in class**



Why the success of DNNs is ~~surprising~~ obvious

- Hierarchical representations are ubiquitous in AI. Computer vision:

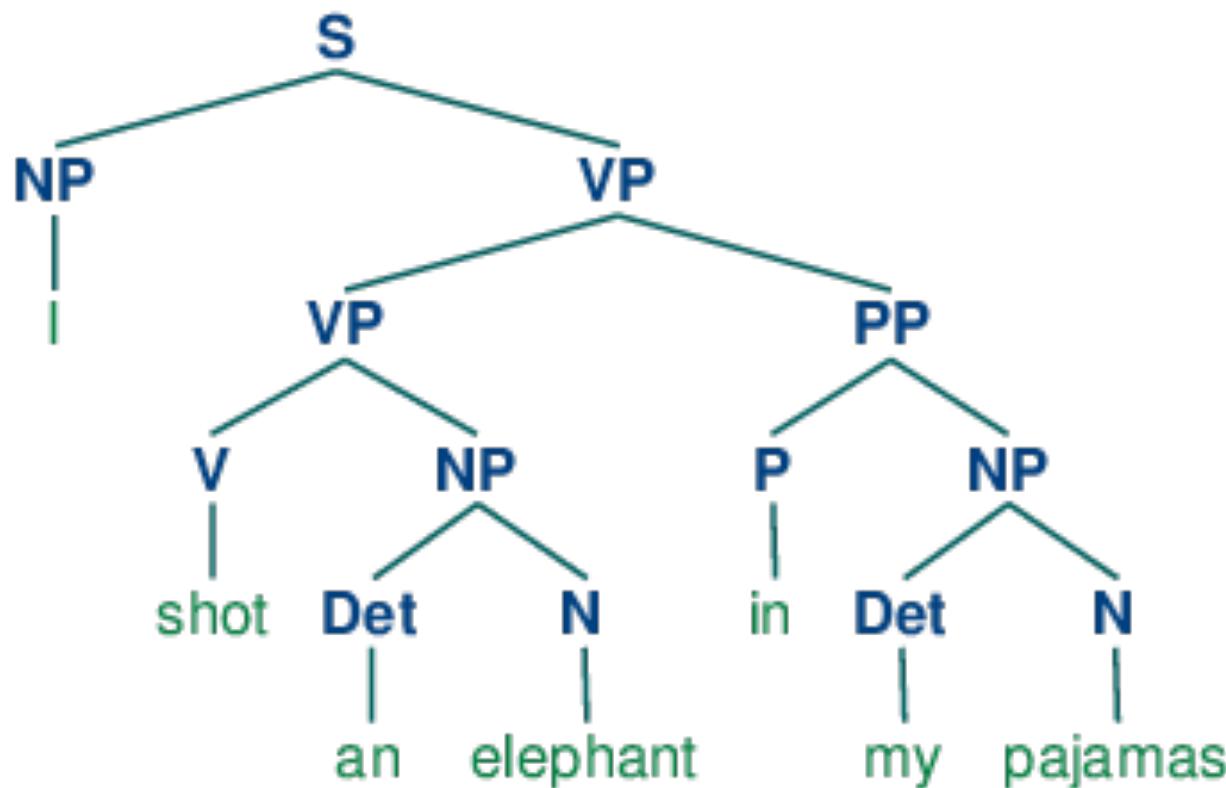


Stages of Visual Representation, David Marr,
1970s

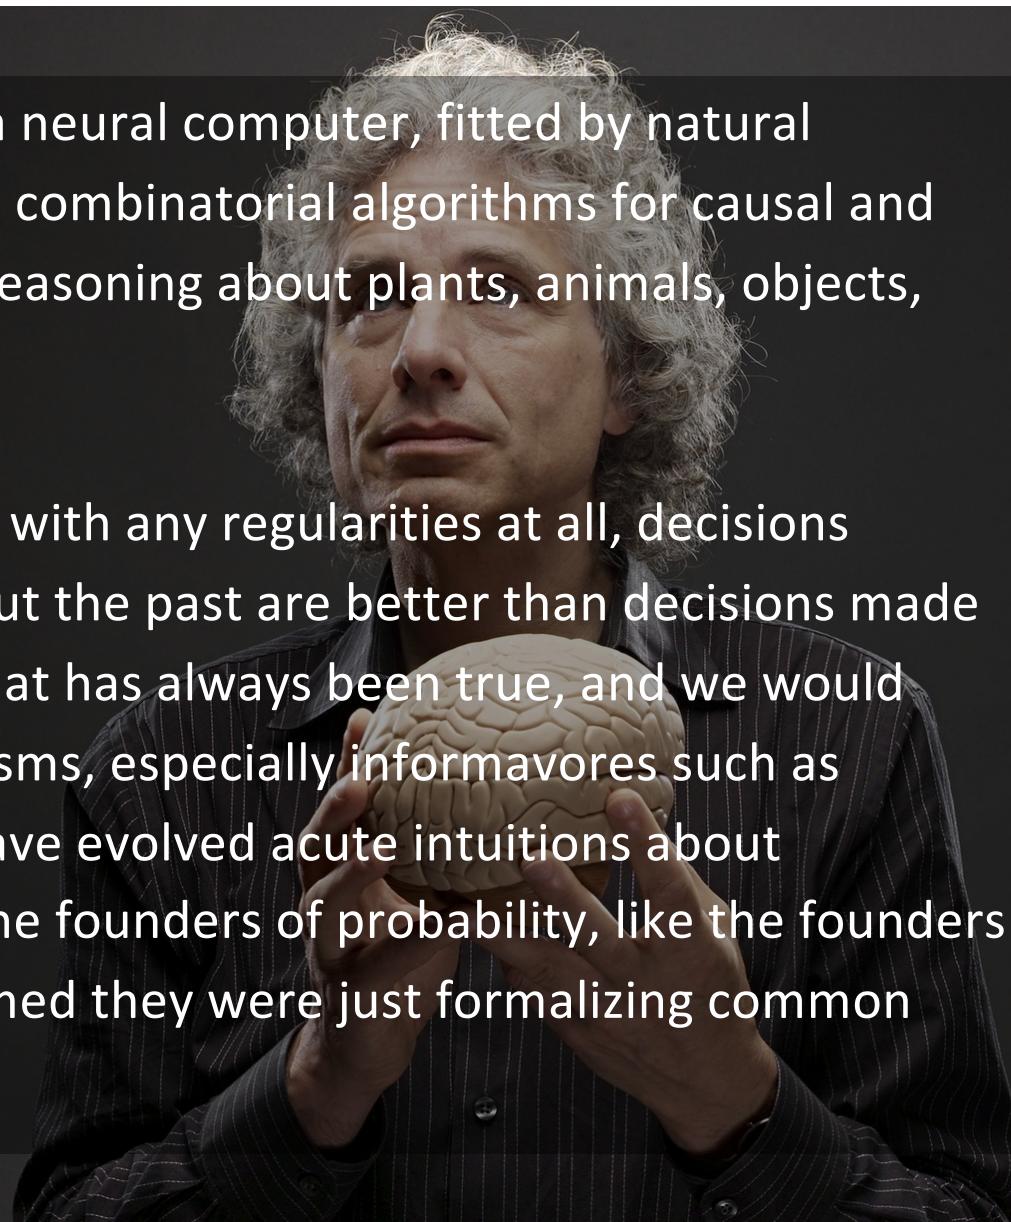


Why the success of DNNs is ~~surprising~~ obvious

- Natural language:



Why the success of DNNs is ~~surprising~~ obvious

A black and white photograph of Steven Pinker, a man with curly grey hair and a beard, wearing a dark pinstripe shirt. He is looking slightly downwards and to his left. In his hands, he holds a large, white, segmented model of a human brain, which is visible in the lower half of the frame.

“The mind is a neural computer, fitted by natural selection with combinatorial algorithms for causal and probabilistic reasoning about plants, animals, objects, and people.”

...

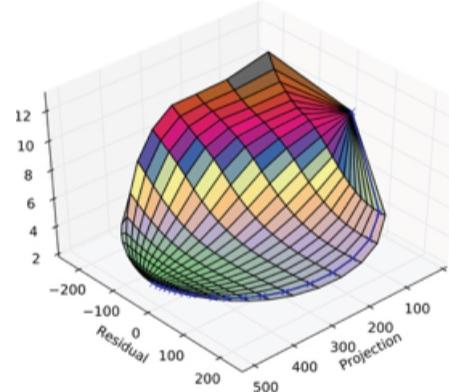
“In a universe with any regularities at all, decisions informed about the past are better than decisions made at random. That has always been true, and we would expect organisms, especially informavores such as humans, to have evolved acute intuitions about probability. The founders of probability, like the founders of logic, assumed they were just formalizing common sense.”

-Steven Pinker, How the Mind Works, 1997, pp. 524, 343.

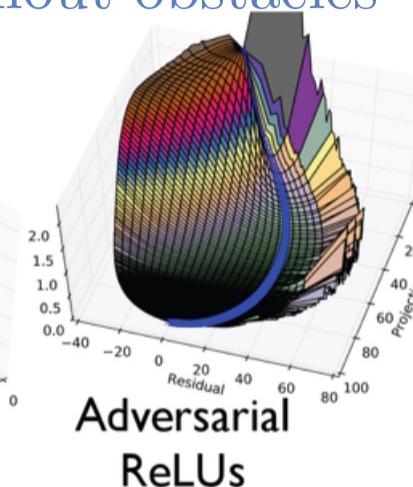
Why the success of DNNs is ~~surprising~~ obvious

- What about greedy optimization?
- Less obvious, but it looks like many learning problems (e.g. image classification) are actually “easy” i.e. have reliable steepest descent paths to a good model.

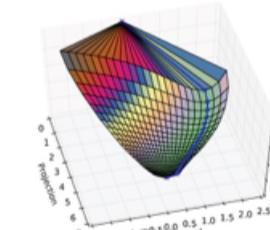
3-D plots without obstacles



LSTM



Adversarial
ReLUs



Factored Linear

Ian Goodfellow – ICLR 2015 Tutorial



What have you learnt?

- Deep Learning is quite successful
- Deep Learning is hyped
- The goal of the course is to review principles and techniques for understanding deep networks.



Have fun!



- Watch the next video and prepare 5 questions! We start with the part on optimizaton. The part on KNN, linear models etc. is only a recap as it was covered in other lectures.
- Also, start to think about your projects? What are interesting topics? What would you enjoy? Prepare a two-pages expose (Latex including clean and complete references)

